

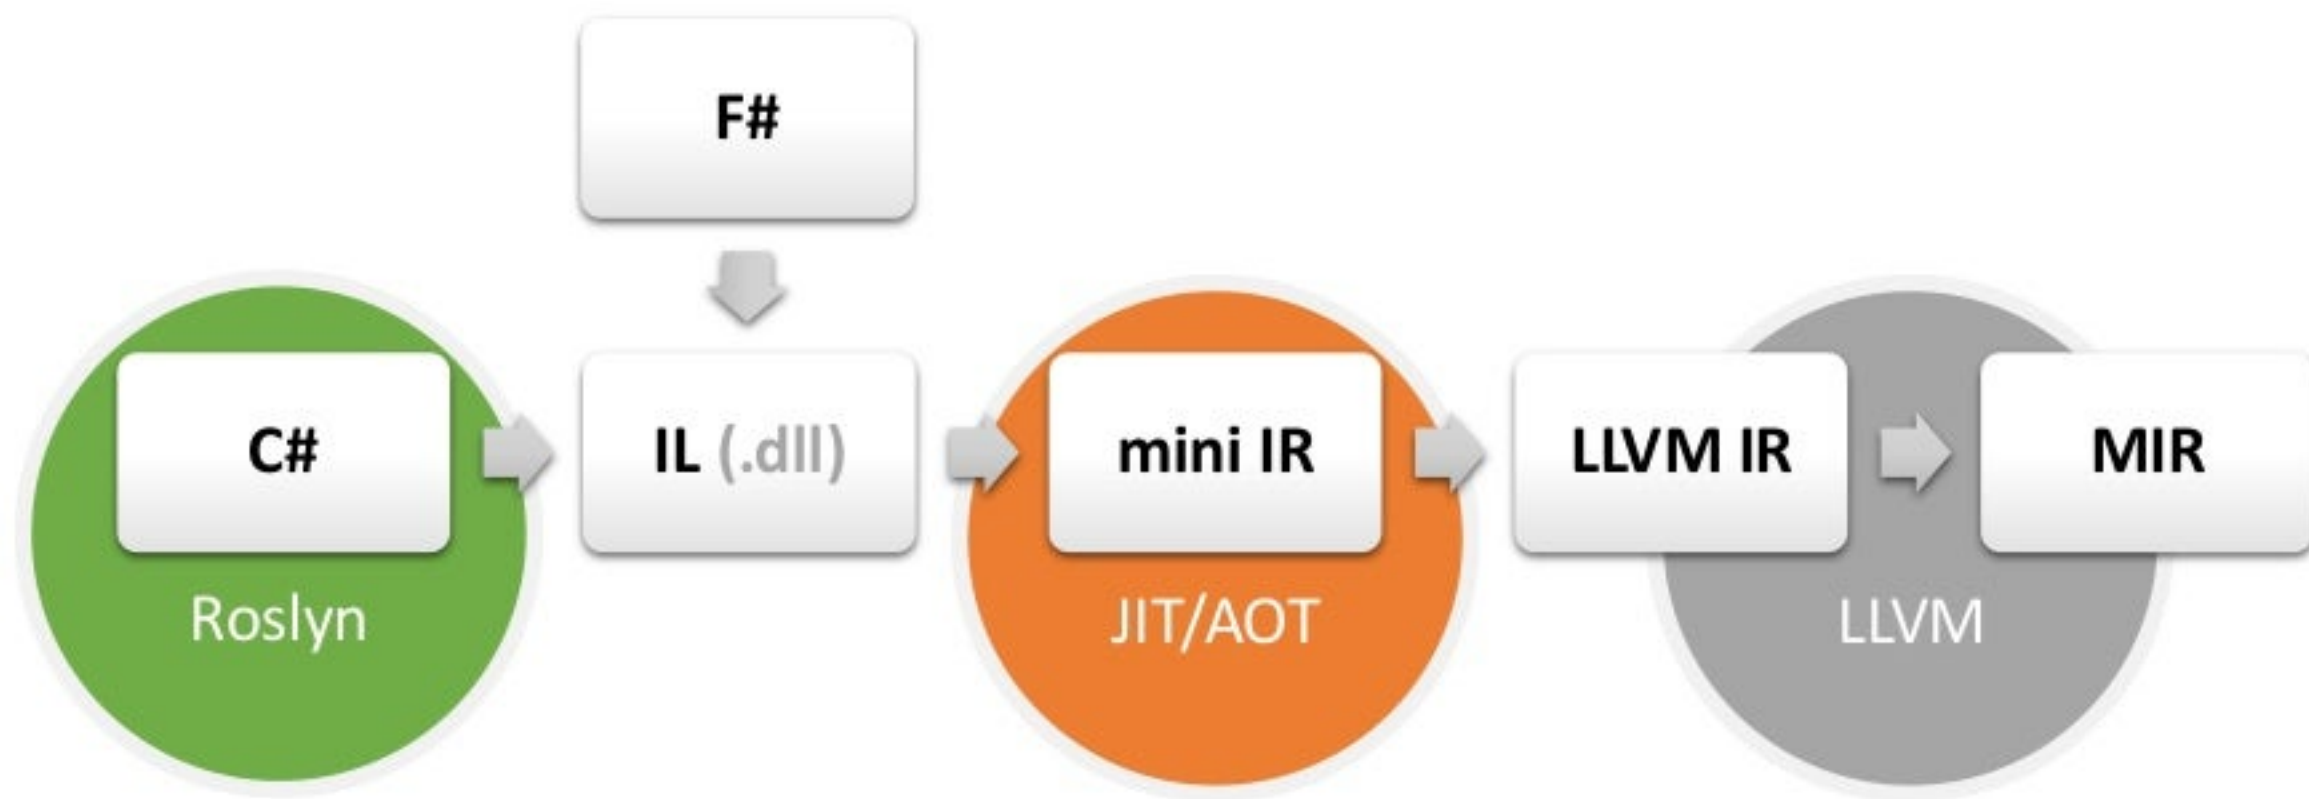
# LLVM back-end for C#



@EgorBo

Senior Software Engineer at Microsoft

# Mono-LLVM: quick intro





# LLVM: C# to MC

```
int MyAdd(int a, int b)
{
    return a + b;
}
```



Roslyn

```
.method private hidebysig static
    int32 MyAdd (
        int32 a,
        int32 b
    ) cil managed
{
    IL_0000: ldarg.0
    IL_0001: ldarg.1
    IL_0002: add
    IL_0003: ret
}
```

# LLVM: C# to MC

```
.method private hidebysig static
    int32 MyAdd (
        int32 a,
        int32 b
    ) cil managed
{
    IL_0000: ldarg.0
    IL_0001: ldarg.1
    IL_0002: add
    IL_0003: ret
}
```



mini

```
define monocc i32 @"Program:MyAdd (int,int)"
    (i32 %arg_a, i32 %arg_b) {
BB0:
    br label %BB3

BB3:
    br label %BB2

BB2:
    %t22 = add i32 %arg_a, %arg_b
    br label %BB1

BB1:
    ret i32 %t22
}
```

# LLVM: C# to MC

```
.method private hidebysig static
    int32 MyAdd (
        int32 a,
        int32 b
    ) cil managed
{
    IL_0000: ldarg.0
    IL_0001: ldarg.1
    IL_0002: add
    IL_0003: ret
}
```



```
define monocc i32 @"Program:MyAdd (int,int)"
    (i32 %arg_a, i32 %arg_b) {
BB0:
    br label %BB3

BB3:                                     (simplifcfg)
    br label %BB2

BB2:
    %t22 = add i32 %arg_a, %arg_b
    br label %BB1

BB1:
    ret i32 %t22
}
```

# LLVM: C# to MC

```
define monocc i32 @"Program:MyAdd (int,int)"(i32 %arg_a, i32 %arg_b) #0
{
BB0:
    %t22 = add i32 %arg_a, %arg_b
    ret i32 %t22
}
```



(I am skipping the machine IR part)

```
lea eax, [rdi + rsi]
ret
```

# LLVM: C# to MC

```
void MySet(byte* array, int len, byte val)
{
    for (int i = 0; i < len; i++)
        array[i] = val;
}
```



mini

```
define monocc void @"MySet"(i8* %arg_array,
                             i32 %arg_len, i32 %arg_val) #0 {
BB0:
    br label %BB3

BB3:
    br label %BB2

BB2:
    br label %BB4

BB4:
    %0 = phi i32 [ 0, %BB2 ], [ %t32, %BB5 ]
    %1 = icmp slt i32 %0, %arg_len
    br i1 %1, label %BB5, label %BB6

BB6:
    br label %BB1

BB5:
    %t23 = sext i32 %0 to i64
    %3 = getelementptr i8, i8* %arg_array, i64 %t23
    %4 = trunc i32 %arg_val to i8
    store i8 %4, i8* %3
    %t32 = add i32 %0, 1
    br label %BB4

BB1:
    ret void
}
```



```

define monocc void @"MySet"(i8* %arg_array,
                           i32 %arg_len, i32 %arg_val) #0 {
BB0:
    br label %BB3

BB3:
    br label %BB2

BB2:
    br label %BB4

BB4:
    %0 = phi i32 [ 0, %BB2 ], [ %t32, %BB5 ]
    %1 = icmp slt i32 %0, %arg_len
    br i1 %1, label %BB5, label %BB6

BB6:
    br label %BB1

BB5:
    %t23 = sext i32 %0 to i64
    %3 = getelementptr i8, i8* %arg_array, i64 %t23
    %4 = trunc i32 %arg_val to i8
    store i8 %4, i8* %3
    %t32 = add i32 %0, 1
    br label %BB4

BB1:
    ret void
}

```

opt

```

define monocc void @"MySet"(i8* %arg_array,
                           i32 %arg_len, i32 %arg_val) #0 {
BB0:
    %0 = icmp sgt i32 %arg_len, 0
    br i1 %0, label %BB5.lr.ph, label %BB1

BB5.lr.ph:
    %1 = trunc i32 %arg_val to i8
    %2 = zext i32 %arg_len to i64
    call void @llvm.memset.p0i8.i64(i8* %arg_array, i8 %1, i64 %2...)
    br label %BB1

BB1:
    ret void
}

```



# LLVM needs hints from front-ends

- Alias-analysis (``noalias``, TBAA, etc)
- PGO-data and `@llvm.expect` (branch-weights)
- Use **GEPs** where possible instead of ``ptrtoint+add+inttoptr``
- Language specific order of optimization passes (maybe even custom passes)
- Alignment hints
- FaultMaps and implicit null-checks
- Use LLVM intrinsics where needed (libcalls, HW intrinsics, etc)
- Don't forget about Fast-Math!
- NIT: Nullability, Escape analysis, etc

# AA is important!

```
A Save/Load + Add new... Vim C x86_64 clang 9.0.0 -g0 -emit-llvm -O2
```

```
1 void MyCopy(int* restrict a, int* b, int len)
2 {
3     for (int i=0; i<len; i++)
4     {
5         a[i] = b[i];
6     }
7 }
8
9
```

```
A 11010 ./a.out .LX0: lib.f: .text // \s+ Intel Demangle Libraries
1 target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
2 target triple = "x86_64-unknown-linux-gnu"
3
4 define dso_local void @MyCopy(i32* noalias nocapture, i32* nocapture readonly, i32) local_unnamed_a
5     %4 = icmp sgt i32 %2, 0
6     br i1 %4, label %5, label %10
7
8 5: ; preds = %3
9     %6 = bitcast i32* %1 to i8*
10    %7 = bitcast i32* %0 to i8*
11    %8 = zext i32 %2 to i64
12    %9 = shl nuw nsw i64 %8, 2
13    call void @llvm.memcpy.p0i8.p0i8.i64(i8* align 4 %7, i8* align 4 %6, i64 %9, i1 false)
14    br label %10
15
16 10: ; preds = %5, %3
17    ret void
18 }
19
```

# LLVM: ffast-math

1

```
float result = (x / 10) + 1 + 1;
```

2

```
float result = (x / 10) + 2;
```

3

```
float result = (x * 0.1) + 2;
```

4

```
float result = fmadd(x, 0.1, 2);
```

CoreCLR, Mono, C++ -O2

```
vdivsd xmm0, xmm0, QWORD PTR .LC0[rip]  
vaddsd xmm0, xmm0, xmm1  
vaddsd xmm0, xmm0, xmm1
```

Mono-LLVM (--llvm --ffast-math)

```
vfmadd132sd xmm0, xmm1, QWORD PTR .LC0[rip]
```



# LLVM: ffast-math

- No NaNs (**NOTE: some apps rely on NaN values, e.g. WPF**)
- No Infs
- No Signed Zeros (**e.g. `Math.Max(-0.0, 0.0)` will return `-0.0`**)
- Approximate functions
- Recognize FMA patterns (**`a * b + c`**)
- Reassociation transformations (**`x + c1 + c2 => x + (c1 + c2)`**)

# C# Math/MathF are LLVM intrinsics!

```
static float Foo(float x)
{
    return MathF.Sqrt(x) * MathF.Sqrt(x);
}
```



Emit  
LLVM IR

```
define float @"Foo" (float %arg_x) {
BB0:
    br label %BB3

BB3:
    br label %BB2

BB2:
    %t19 = call fast float @llvm.sqrt.f32(float %arg_x)
    %t21 = call fast float @llvm.sqrt.f32(float %arg_x)
    %t22 = fmul fast float %t19, %t21
    br label %BB1

BB1:
    ret float %t22
}

declare float @llvm.sqrt.f32(float)
```

# C# Math/MathF are LLVM intrinsics!

```
define float @"Foo" (float %arg_x) {
```

```
BB0:  
  br label %BB3
```

```
BB3:  
  br label %BB2
```

```
BB2:  
  %t19 = call fast float @llvm.sqrt.f32(float %arg_x)  
  %t21 = call fast float @llvm.sqrt.f32(float %arg_x)  
  %t22 = fmul fast float %t19, %t21  
  br label %BB1
```

```
BB1:  
  ret float %t22  
}
```

```
declare float @llvm.sqrt.f32(float)
```



opt

```
define float @"Foo" (float %arg_x) {
```

```
BB0:  
  ret float %arg_x  
}
```



# C# Math/MathF are LLVM intrinsics!

`Math.Abs(X) * Math.Abs(X)`  $\Rightarrow$  `X * X`

`Math.Sin(X) / Math.Cos(X)`  $\Rightarrow$  `Math.Tan(X)`

`Math.Sqrt(X) * Math.Sqrt(X)`  $\Rightarrow$  `X`

`Math.Sin(-X)`  $\Rightarrow$  `-Math.Sin(X)`

`MathF.Pow(X, 0.5f)`  $\Rightarrow$  `MathF.Sqrt(X)`

`MathF.Pow(X, 2)`  $\Rightarrow$  `X * X`

`MathF.Pow(X, 4)`  $\Rightarrow$  `(X * X) * (X * X)`

# Significant boost in some FP benchmarks:

Benchmark	Fast-Math	No Fast-Math
Benchmark.BenchF.BenchMk2.Test	1.00	2.25
Benchmark.BenchF.BenchMrk.Test	1.00	2.37
Benchmark.BenchF.Lorenz.Test	1.00	1.85
Burgers.Test0	1.00	1.39
Burgers.Test1	1.00	2.30
Burgers.Test2	1.00	2.32
System.MathBenchmarks.Double.Cosh	1.00	1.50

(CPU without FMA)

# One of those benchmarks (Lorenz equations):

```

    x_arg = s_x + hdiv2 * k2;
    y_arg = s_y + hdiv2 * l2;
    z_arg = s_z + hdiv2 * m2;

    k3 = F(t_arg, x_arg, y_arg, z_arg);
    l3 = G(t_arg, x_arg, y_arg, z_arg);
    m3 = H(t_arg, x_arg, y_arg, z_arg);

    t_arg = s_t + s_h;
    x_arg = s_x + s_h * k3;
    y_arg = s_y + s_h * l3;
    z_arg = s_z + s_h * m3;

    k4 = F(t_arg, x_arg, y_arg, z_arg);
    l4 = G(t_arg, x_arg, y_arg, z_arg);
    m4 = H(t_arg, x_arg, y_arg, z_arg);

    s_x = s_x + hdiv6 * (k1 + 2.0 * k2 + 2.0 * k3 + k4);
    s_y = s_y + hdiv6 * (l1 + 2.0 * l2 + 2.0 * l3 + l4);
    s_z = s_z + hdiv6 * (m1 + 2.0 * m2 + 2.0 * m3 + m4);
    s_t = t_arg;
}

return true;
}

private static double F(double t, double x, double y, double z)
{
    return (10.0 * (y - x));
}

private static double G(double t, double x, double y, double z)
{
    return (x * (28.0 - z) - y);
}

private static double H(double t, double x, double y, double z)
{
    return (x * y - (8.0 * z) / 3.0);
}
```



Don't forget to feed the **LLVM!**  
with some **PGO** data for better codegen

```
static int Max(int x, int y)
{
    return x > y ? x : y;
}
```

### Default

```
define i32 @Max(i32 %0, i32 %1) {
    %3 = icmp sgt i32 %0, %1
    %4 = select i1 %3, i32 %0, i32 %1
    ret i32 %4
}
```

```
mov eax, esi
cmp edi, esi
cmovge eax, edi
ret
```

### With profile data (PGO)

```
define i32 @Max(i32 %0, i32 %1) {
    %3 = icmp sgt i32 %0, %1 !prof 0
    %4 = select i1 %3, i32 %0, i32 %1
    ret i32 %4
}
```

```
!0 = !{"branch_weights", i32 1000, i32 1}
```

```
mov eax, edi
cmp edi, esi
jle .LBB1_1
ret
.LBB1_1:
mov eax, esi
ret
```

LLVM IR source #1 X

X

opt (trunk) (Editor #1, Compiler #2) LLVM IR X

A Save/Load + Add new... Vim

LLVM IR

```
1 define i32 @Max(i32, i32) {  
2   %3 = icmp sgt i32 %0, %1  
3   %4 = select i1 %3, i32 %0, i32 %1  
4   ret i32 %4  
5 }
```

opt (trunk)

-pgo-instr-gen -instrprof

A 11010 .a.out .LX0: lib.f: .text // \s+ Intel Demangle Libraries + Add new... Add tool...

```
1 define i32 @Max(i32 %0, i32 %1) {  
2   %pgocount = load i64, i64* @getelementptr inbounds ([2 x i64], [2 x i64]* @__profc_Max, i64 0, i64 0)  
3   %3 = add i64 %pgocount, 1  
4   store i64 %3, i64* @getelementptr inbounds ([2 x i64], [2 x i64]* @__profc_Max, i64 0, i64 0)  
5   %4 = icmp sgt i32 %0, %1  
6   %5 = zext i1 %4 to i64  
7   %pgocount1 = load i64, i64* @getelementptr inbounds ([2 x i64], [2 x i64]* @__profc_Max, i64 0, i64 1)  
8   %6 = add i64 %pgocount1, %5  
9   store i64 %6, i64* @getelementptr inbounds ([2 x i64], [2 x i64]* @__profc_Max, i64 0, i64 1)  
10  %7 = select i1 %4, i32 %0, i32 %1  
11  ret i32 %7  
12 }  
13  
14 declare void @llvm.instrprof.increment(i8*, i64, i32, i32) #0  
15  
16 declare void @llvm.instrprof.increment.step(i8*, i64, i32, i32, i64) #0  
17  
18 define linkonce_odr hidden i32 @__llvm_profile_runtime_user() #1 comdat {  
19   %1 = load i32, i32* @__llvm_profile_runtime  
20   ret i32 %1  
21 }  
22  
23 define internal void @__llvm_profile_register_functions() unnamed_addr {  
24   call void @__llvm_profile_register_function(i8* bitcast ([ i64, i64, i64*, i8*, i8*, i32, [2 x i16] ]* @__profd_Max to i8*))  
25   call void @__llvm_profile_register_names_function(i8* @getelementptr inbounds ([13 x i8], [13 x i8]* @__llvm_prf_nm, i32 0, i32 0), i64  
26   ret void  
27 }  
28  
29 declare void @__llvm_profile_register_function(i8*)  
30  
31 declare void @__llvm_profile_register_names_function(i8*, i64)  
32  
33 define internal void @__llvm_profile_init() unnamed_addr #1 {  
34   call void @__llvm_profile_register_functions()  
35   ret void  
36 }  
37  
38 attributes #0 = { nounwind }  
39 attributes #1 = { noinline }
```



# PGO and switch

```
char* format(char format)
{
    switch (format)
    {
        case 'G':
            return defaultFormat();

        case 'X':
            return hexFormat();

        case 'f':
            return floatFormat();

        case 'p':
            return percentFormat();
    }
    return defaultFormat();
}
```

```
format(int): # @format(int)
    cmp edi, 101
    jg .LBB0_4
    cmp edi, 71
    je .LBB0_8
    cmp edi, 88
    jne .LBB0_8
    jmp hexFormat() # TAILCALL
.LBB0_4:
    cmp edi, 102
    je .LBB0_9
    cmp edi, 112
    jne .LBB0_8
    jmp percentFormat() # TAILCALL
.LBB0_8:
    jmp defaultFormat() # TAILCALL
.LBB0_9:
    jmp floatFormat() # TAILCALL
```

# PGO and switch

```
char* format(char format)
{
    switch (__builtin_expect(format, 'X'))
    {
        case 'G':
            return defaultFormat();

        case 'X':
            return hexFormat();

        case 'f':
            return floatFormat();

        case 'p':
            return percentFormat();
    }
    return defaultFormat();
}
```

```
format(int): # @format(int)
    movsxd rax, edi
    cmp rax, 88
    jne .LBB0_2
    jmp hexFormat() # TAILCALL
.LBB0_2:
    cmp rax, 112
    je .LBB0_6
    cmp rax, 102
    je .LBB0_7
    cmp rax, 71
    jmp defaultFormat() # TAILCALL
.LBB0_6:
    jmp percentFormat() # TAILCALL
.LBB0_7:
    jmp floatFormat() # TAILCALL
```

# PGO and Guarded Devirtualization

```
static void Foo(IAnimal animal)
{
    animal.MakeSound();
}
```



```
static void Foo(IAnimal animal)
{
    if (animal is Dog dog)
        dog.Bark();
    else
        animal.MakeSound();
}
```

# Mono-LLVM

- We use [github.com/dotnet/llvm-project](https://github.com/dotnet/llvm-project) fork
- Currently target LLVM 6 but on our way to LLVM 9
- We currently use LLVM for:
  - AOT (opt + llc)
  - JIT (`legacy::PassManager`, ORCv1)



# Mono-LLVM

- `opt -O2` is only for C++
- We have a lot of additional checks: null-checks, bound-checks
- We have to insert safe-points (`-place-safepoints`) for non-preemptive mode

```
void MyMethod(byte* array, int len, byte val)
{
    if (unlikely(gcRequested))
        performGC();

    for (int i = 0; i < len; i++)
    {
        if (unlikely(gcRequested))
            performGC();

        array[i] = val;
    }

    if (unlikely(gcRequested))
        performGC();
}
```

# How do we use optimizations?

- **LLVM AOT:**

`opt -O2 -place-safepoints`

***`opt -O2`**: These pass pipelines make a good starting point for an optimizing compiler for any language, but they have been carefully tuned for C and C++, not your target language.*

- **LLVM JIT:**

**PassManager** with a small subset of the most useful passes in a specific order:

- simplifycfg
- sroa
- lower-expect
- instcombine
- licm
- simplifycfg
- lcssa
- indvars
- loop-deletion
- gvn
- memcpyopt
- sccp
- bdce
- instcombine
- dse
- simplifycfg

```
gc
mono_aot_HelloWorld_icall_cold_wrapper_265
llvm_code_start
mono_aot_HelloWorldinit_method
mono_aot_HelloWorldinit_method_gshared_mrgctx
mono_aot_HelloWorldinit_method_gshared_this
mono_aot_HelloWorldinit_method_gshared_vtable
PP_Main
TimeSpanTokenizer_ctor_System_ReadOnlySpan_1_char
TimeSpanTokenizer_ctor_System_ReadOnlySpan_1_char_int
TimeSpanTokenizer_get_NextChar
```

Module Verifier -verify

Instrument function entry/exit with calls to e.g. mcount() (pre inlining) -??

Simplify the CFG -simplifycfg

SROA -sroa

Early CSE -early-cse

Lower 'expect' Ininsics -lower-expect

Promote Memory to Register -mem2reg

Combine redundant instructions -instcombine

Simplify the CFG -simplifycfg

Remove unused exception handling info -prune-eh

Function Integration/Inlining -inline

Deduce function attributes -functionattrs

Remove unused exception handling info -prune-eh

Function Integration/Inlining -inline

Deduce function attributes -functionattrs

SROA -sroa

Early CSE w/ MemorySSA -memoryssa

Speculatively execute instructions if target has divergent branches -??

Jump Threading -jump-threading

Value Propagation -correlated-propagation

Simplify the CFG -simplifycfg

Combine redundant instructions -instcombine

Conditionally eliminate dead library calls -??

PGOMemOpsize -??

Tail Call Elimination -tailcallelim

Simplify the CFG -simplifycfg

Reassociate expressions

Canonicalize natural loops -loop-simplify

LCSSA Verifier -??

Loop-Closed SSA Form Pass -lcssa

Simplify the CFG -simplifycfg

Combine redundant instructions -instcombine

Canonicalize natural loops -loop-simplify

LCSSA Verifier -??

Loop-Closed SSA Form Pass -lcssa

MergedLoadStoreMotion -??

Global Value Numbering -gvn

MemCpy Optimization -memcpyopt

Sparse Conditional Constant Propagation -sccp

Demand bits analysis -??

Bit-Tracking Dead Code Elimination -bdce

Combine redundant instructions -instcombine

Jump Threading -jump-threading

Value Propagation -correlated-propagation

Dead Store Elimination -dse

```
: Function Attrs: noline uwtable
define hidden moncc @TimeSpanTokenizer_get_NextChar(i64* %this) #4 {
BB0:
    %0 = icmp eq i64* %this, null
    br i1 %0, label %EX_BB3, label %NOEX_BB4, !make.implicit !2
EX_BB3:
    ; preds = %BB0
    call void @p_1_plt__it_icall_llvm_throw_corlib_exception_abs_trampoline_llvm(i32 204)
    unreachable
NOEX_BB4:
    ; preds = %BB0
    %1 = getelementptr i64, i64* %this, i64 2
    %2 = bitcast i64* %1 to i32*
    %t22 = load i32, i32* %2, align 4
    %t19 = add i32 %t22, 1
    store i32 %t19, i32* %2, align 4
    %3 = getelementptr i64, i64* %this, i64 1
    %4 = bitcast i64* %3 to i32*
    %t30 = load i32, i32* %4, align 4, range !17
    %5 = icmp ult i32 %t19, %t30
    br i1 %5, label %BB4, label %BB1
BB4:
    ; preds = %NOEX_BB4
    %t42 = sext i32 %t19 to i64
    %6 = sext i32 %t30 to i64
    %7 = icmp ugt i64 %6, %t42
    br i1 %7, label %NOEX_BB13, label %EX_BB12
EX_BB12:
    ; preds = %BB4
    call void @p_1_plt__it_icall_llvm_throw_corlib_exception_abs_trampoline_llvm(i32 166)
    unreachable
NOEX_BB13:
    ; preds = %BB4
    %t35 = load i64, i64* %this, align 8
    %t36 = shl nsw i64 %t42, 1
    %t37 = add i64 %t35, %t36
    %8 = inttoptr i64 %t37 to i16*
    %9 = load i16, i16* %8, align 2
    br label %BB1
BB1:
    ; preds = %NOEX_BB4, %NOEX_BB13
    %10 = phi i16 [ %9, %NOEX_BB13 ], [ 0, %NOEX_BB4 ]
    ret i16 %10
}
```

```
: Function Attrs: noline uwtable
define hidden moncc @TimeSpanTokenizer_get_NextChar(i64* %this) #4 {
BB0:
    %0 = icmp eq i64* %this, null
    br i1 %0, label %EX_BB3, label %NOEX_BB4, !make.implicit !2
EX_BB3:
    ; preds = %BB0
    call void @p_1_plt__it_icall_llvm_throw_corlib_exception_abs_trampoline_llvm(i32 204)
    unreachable
NOEX_BB4:
    ; preds = %BB0
    %1 = getelementptr i64, i64* %this, i64 2
    %2 = bitcast i64* %1 to i32*
    %t22 = load i32, i32* %2, align 4
    %t19 = add i32 %t22, 1
    store i32 %t19, i32* %2, align 4
    %3 = getelementptr i64, i64* %this, i64 1
    %4 = bitcast i64* %3 to i32*
    %t30 = load i32, i32* %4, align 4, range !17
    %5 = icmp ult i32 %t19, %t30
    br i1 %5, label %BB4, label %BB1
BB4:
    ; preds = %NOEX_BB4
    %t42 = sext i32 %t19 to i64
    %6 = sext i32 %t30 to i64
    %7 = icmp ugt i64 %6, %t42
    br i1 %7, label %NOEX_BB13, label %EX_BB12
EX_BB12:
    ; preds = %BB4
    call void @p_1_plt__it_icall_llvm_throw_corlib_exception_abs_trampoline_llvm(i32 166)
    unreachable
NOEX_BB13:
    ; preds = %BB4
    %t35 = load i64, i64* %this, align 8
    %t36 = mul i64 %t42, 2
    %t37 = add i64 %t35, %t36
    %8 = inttoptr i64 %t37 to i16*
    %9 = load i16, i16* %8, align 2
    br label %BB1
BB1:
    ; preds = %NOEX_BB4, %NOEX_BB13
    %10 = phi i16 [ %9, %NOEX_BB13 ], [ 0, %NOEX_BB4 ]
    ret i16 %10
}
```



# LLVM: tons of optimizations

- `-adce`: Aggressive Dead Code Elimination
- `-always-inline`: Inliner for `always_inline` functions
- `-argpromotion`: Promote 'by reference' arguments to scalars
- `-bb-vectorize`: Basic-Block Vectorization
- `-block-placement`: Profile Guided Basic Block Placement
- `-break-crit-edges`: Break critical edges in CFG
- `-codegenprepare`: Optimize for code generation
- `-constmerge`: Merge Duplicate Global Constants
- `-constprop`: Simple constant propagation
- `-dce`: Dead Code Elimination
- `-deadargelim`: Dead Argument Elimination
- `-deadtypeelim`: Dead Type Elimination
- `-die`: Dead Instruction Elimination
- `-dse`: Dead Store Elimination
- `-functionattrs`: Deduce function attributes
- `-globaldce`: Dead Global Elimination
- `-globalopt`: Global Variable Optimizer
- `-gvn`: Global Value Numbering
- `-indvars`: Canonicalize Induction Variables
- `-inline`: Function Integration/Inlining
- `-instcombine`: Combine redundant instructions
- `-aggressive-instcombine`: Combine expression patterns
- `-internalize`: Internalize Global Symbols
- `-ipconstprop`: Interprocedural constant propagation
- `-ipsccp`: Interprocedural Sparse Conditional Constant Propagation
- `-jump-threading`: Jump Threading
- `-lcssa`: Loop-Closed SSA Form Pass
- `-licm`: Loop Invariant Code Motion
- `-loop-deletion`: Delete dead loops
- `-loop-extract`: Extract loops into new functions
- `-loop-extract-single`: Extract at most one loop into a new function
- `-loop-reduce`: Loop Strength Reduction
- `-loop-rotate`: Rotate Loops
- `-loop-simplify`: Canonicalize natural loops
- `-loop-unroll`: Unroll loops
- `-loop-unroll-and-jam`: Unroll and Jam loops
- `-loop-unswitch`: Unswitch loops
- `-loweratomic`: Lower atomic intrinsics to non-atomic form
- `-lowerinvoke`: Lower invokes to calls, for unwindless code generators
- `-lowerswitch`: Lower SwitchInsts to branches
- `-mem2reg`: Promote Memory to Register
- `-mencpyopt`: MemCpy Optimization
- `-mergefunc`: Merge Functions
- `-mergereturn`: Unify function exit nodes
- `-partial-inliner`: Partial Inliner
- `-prune-eh`: Remove unused exception handling info
- `-reassociate`: Reassociate expressions
- `-reg2mem`: Demote all values to stack slots
- `-sroa`: Scalar Replacement of Aggregates
- `-sccp`: Sparse Conditional Constant Propagation
- `-simplifycfg`: Simplify the CFG
- `-sink`: Code sinking
- `-strip`: Strip all symbols from a module
- `-strip-dead-debug-info`: Strip debug info for unused symbols
- `-strip-dead-prototypes`: Strip Unused Function Prototypes
- `-strip-debug-declare`: Strip all `llvm.dbg.declare` intrinsics
- `-strip-nondebug`: Strip all symbols, except dbg symbols, from a module
- `-tailcallelim`: Tail Call Elimination



# Null-checks, Bound-checks

```
static int Test(int[] array)
{
    return array[42];
}
```

# Null-checks, Bound-checks

```
static int Test(int[] array)
{
    if (array == null) ThrowNRE();
    return array[42];
}
```

# Null-checks, Bound-checks

```
static int Test(int[] array)
{
    if (array == null) ThrowNRE();
    if ((uint)array.Length <= 42) ThrowOBE();
    return array[42];
}
```

# Null-checks, Bound-checks

```
static int Test(int[] array)
{
    if (array == null) ThrowNRE();
    if ((uint)array.Length <= 42) ThrowOBE();
    return array[42];
}
```



[LLVM Home](#) | [Documentation](#) » [Reference](#) »

FaultMaps and implicit checks

```
push    rax
cmp     DWORD PTR [rdi+0x18],0x0
je      1d <Test__int____+0x1d>
mov     eax,DWORD PTR [rdi+0x20]
pop     rcx
ret
```

```
movabs  rax,0x1f1f030
mov     edi,0xcc
call    QWORD PTR [rax]
movabs  rax,0x1f1f040
mov     edi,0xa6
call    QWORD PTR [rax]
```



# InductiveRangeCheckElimination Pass

```
public static void Zero1000Elements(int[] array)
{
    for (int i = 0; i < 1000; i++)
        array[i] = 0; // bound checks will be inserted here
}
```

*"If your language uses range checks, consider using the IRCE pass. It is not currently part of the standard pass order."*

# InductiveRangeCheckElimination Pass

```
public static void Zero1000Elements(int[] array)
{
    int limit = Math.Min(array.Length, 1000);

    for (int i = 0; i < limit; i++)
        array[i] = 0; // bound checks are not needed here!

    for (int i = limit; i < 1000; i++)
        array[i] = 0; // bound checks are needed here

    // so at least we could "zero" first `limit` elements without bound checks
}
```

# InductiveRangeCheckElimination Pass

```
public static void Zero1000Elements(int[] array)
{
    int limit = Math.Min(array.Length, 1000);

    for (int i = 0; i < limit - 3; i += 4)
    {
        array[i] = 0;
        array[i+1] = 0;
        array[i+2] = 0;
        array[i+3] = 0;
    }

    for (int i = limit; i < 1000; i++)
        array[i] = 0; // bound checks are needed here

    // so at least we could "zero" first `limit` elements without bound checks
}
```

Now we can even unroll the first loop!

# System.Runtime.Intrinsics.\*

```
static uint Foo(uint x)
{
    return Lzcnt.LeadngZeroCount(x);
}
```



```
define i32 @Foo(i32 %x)
{
    %0 = call i32 @llvmctlz.i32(i32 %x, i1 true)
    ret i32 %0
}
```



```
; x86 with lzcnt
lzcnt eax, edi
ret
```

```
; Arm64 (AArch64)
cls w0, w0
ret
```





# What's wrong with this benchmark?

```
[Benchmark]
public bool IsNaN(float value)
{
    bool result = false;

    for (int i = 0; i < 1000000; i++)
    {
        result &= float.IsNaN(value);
        value += 1.0f;
    }

    return result;
}
```



```
[Benchmark]
public bool IsNaN(float value)
{
    return false;
}
```

# Pros & Cons

- **Pros:**

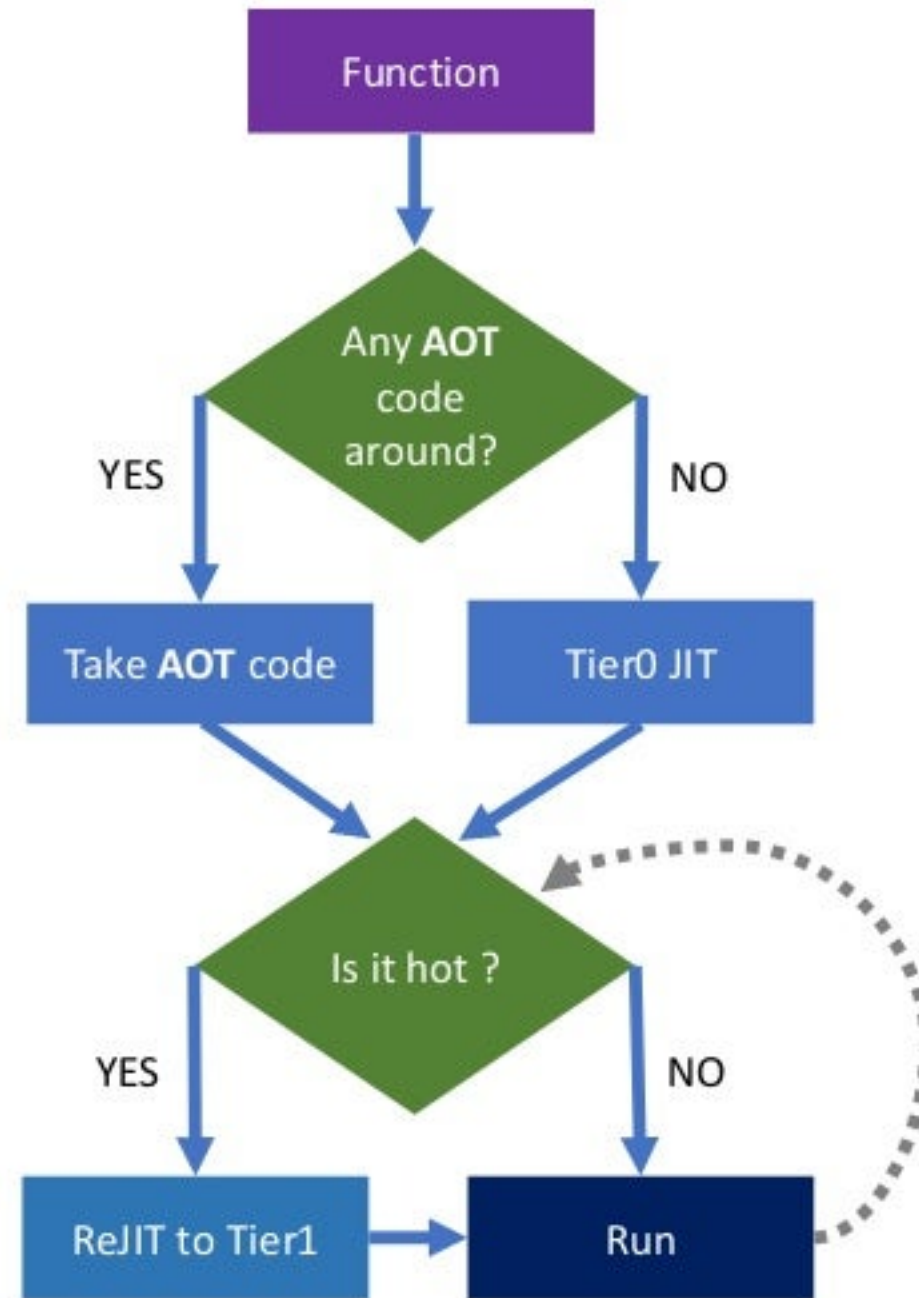
- Tons of optimizations, frontends aren't required to emit optimized IR.
- High MC quality and performance, especially for AArch64!
- A lot of smart people and huge companies invest into LLVM:
  - Apple, Google, Samsung, Intel, Microsoft, etc.

- **Cons:**

- Deadly slow to compile and optimize stuff. LLVM JIT is an oxymoron (unless it's used in a tiered compilation)
- LLVM JIT adds 30-50mb to your application
- Not super friendly for managed languages (Precise GC, "movable" pointers, etc)

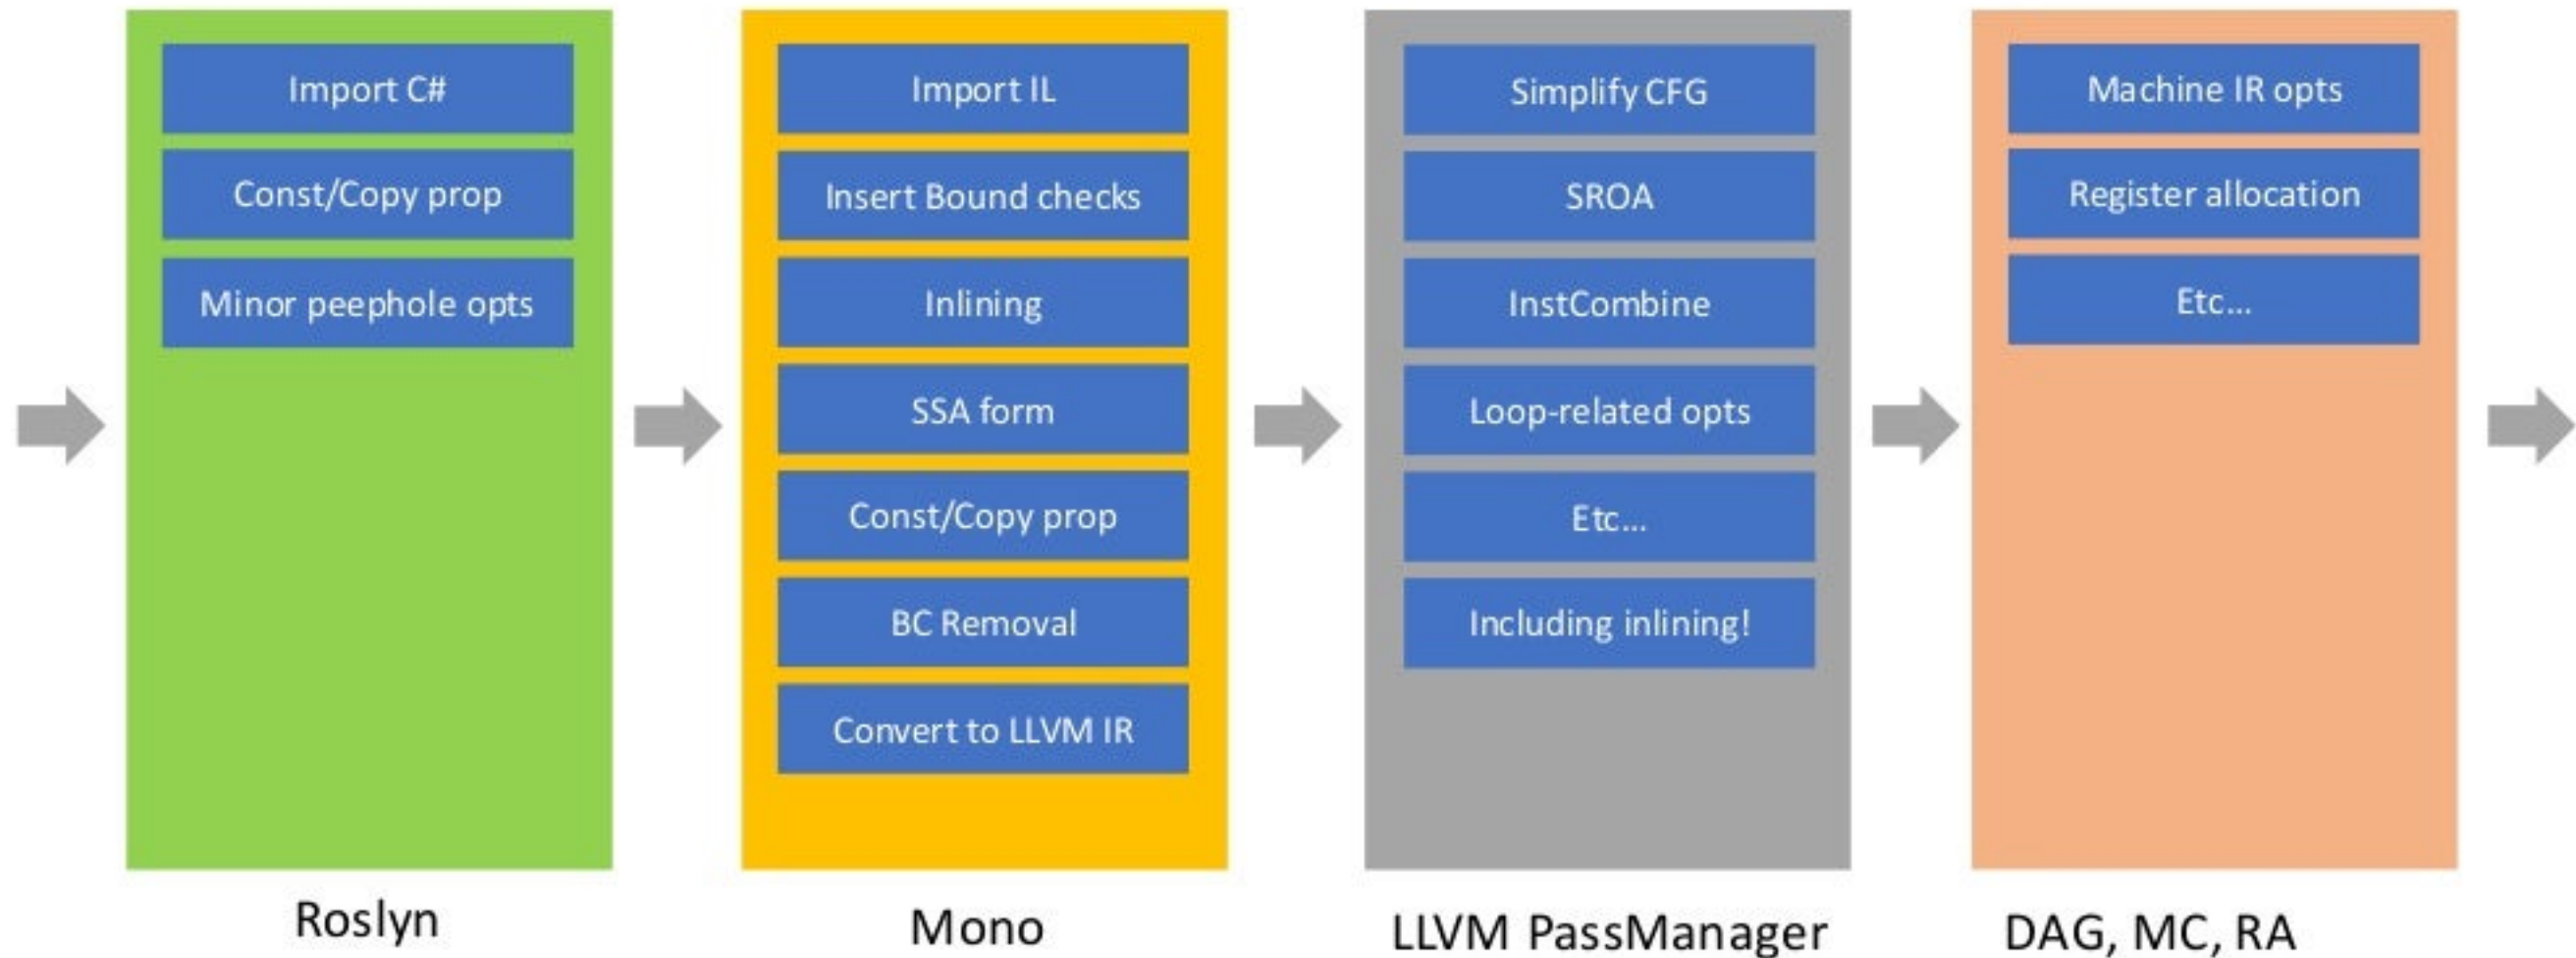
# Tiered JIT

40





# Optimizations... optimizations everywhere!



# Runtimes for C# (IL):

- **.NET Framework 4.x**
  - JIT
  - AOT (Fragile NGEN)
- **CoreCLR**
  - JIT (Tiered)
  - AOT (ReadyToRun – R2R)
- **Mono**
  - JIT
  - (Full)AOT
  - LLVM (AOT, FullAOT, JIT)
  - Interpreter
- **CoreRT**
- Unity **IL2CPP**
- Unity **Burst**

# Q&A

twitter: [EgorBo](#)

blog: [EgorBo.com](#)



EgorBo

Senior Software Engineer at Microsoft